

S3 bucket: Simple Storage Service



A storage space where any type of files can be stored and accessed

Creation

- Bucket name must be unique
- A Bucket can not be transferable across accounts
- A Bucket can be cloned from other bucket

Versioning

- By enabling versioning multiple versions of same doc can be maintained
- Most recent doc will be accessed from any services

Server Access Logging:

- To track the logs related to this bucket
- Store the logs in an another bucket

Tags

- Cost allocation Tags are used to tag the buckets to track the cost of the buckets
- Tags are KEY VALUES pairs
- AWS will generate a 'createdBy' tag by default if user doesn't create any tag after creation of the bucket

S3storage Class

These are the storage classes for amazon s3

- ***standard*** Frequently accessed data
- ***Intelligent-Tiering***
- This is used for objects access is unpredictable
- This will have two storage tiers One is optimized for frequently access and one for the infrequent the pricing is different for these tiers
- Based on the objects access patterns object will be between the two tiers
- ***Standard IA (infrequent access)***
- For Objects used less frequently but expects rapid access when needed this is used ideally for backups
- ***Standard One Zone-IA***
- It is same as Standard IA except the standard IA will store the object across 3 zones but one-zone IA will be stored in only one zone this is suitable for secondary backups
- ***Glacier***
- This is used for archival purpose public access is not possible for the objects under glacier
- Glacier provides three retrieval options that range from

a few minutes to hours

- ***Glacier Deep Archive***
- lowest-cost storage class and supports long-term retention and digital preservation for data that may be accessed once or twice in a year.
- This is ideal solution for data